

Class 08 Lab

Renee Zuhars (PID: A17329856)

Table of contents

1. Exploratory Data Analysis	2
Question 1	5
Question 2	5
Question 3	5
2. Principal Component Analysis	6
Question 4	8
Question 5	8
Question 6	8
Question 7	9
Question 8	10
Question 9	14
Question 10	15
3. Hierarchical Clustering	15
Question 11	15
Question 12	16
Question 13	17
4. K-means clustering (skipped)	19
Question 14 (skipped)	19
5. Combining Methods	19
Question 15	23
Question 16	24
6. Sensitivity / Specificity	24
Question 17 (skipped, running out of time)	25
7. Prediction	25
Question 18	26

1. Exploratory Data Analysis

```
#First, save the input data file into the Project Directory
fna.data <- read.csv("WisconsinCancer.csv")

# Input the data and store as wisc.df
wisc.df <- data.frame(fna.data, row.names=1)

head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217

	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41

84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			
84348301	0.17300			
84358402	0.07678			
843786	0.12440			

```
# Remove the first column (diagnosis) from the data.frame
wisc.data <- wisc.df[,-1]

# View the dataset
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780
	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
842302	0.27760	0.3001	0.14710	0.2419	
842517	0.07864	0.0869	0.07017	0.1812	

84300903	0.15990	0.1974	0.12790	0.2069
84348301	0.28390	0.2414	0.10520	0.2597
84358402	0.13280	0.1980	0.10430	0.1809
843786	0.17000	0.1578	0.08089	0.2087
	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.07871	1.0950	0.9053	8.589
842517	0.05667	0.5435	0.7339	3.398
84300903	0.05999	0.7456	0.7869	4.585
84348301	0.09744	0.4956	1.1560	3.445
84358402	0.05883	0.7572	0.7813	5.438
843786	0.07613	0.3345	0.8902	2.217
	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	0.006399	0.04904	0.05373	0.01587
842517	0.005225	0.01308	0.01860	0.01340
84300903	0.006150	0.04006	0.03832	0.02058
84348301	0.009110	0.07458	0.05661	0.01867
84358402	0.011490	0.02461	0.05688	0.01885
843786	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			

```
84348301          0.17300
84358402          0.07678
843786           0.12440
```

```
# Create diagnosis vector for later
diagnosis <- factor(wisc.df$diagnosis)
```

Question 1

How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations in this dataset.

Question 2

How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

There are 212 malignant diagnoses.

Question 3

How many variables/features in the data are suffixed with `_mean`?

```
sum(grepl("_mean$", names(wisc.data)))
```

```
[1] 10
```

There are 10 features that are suffixed with `_mean`.

2. Principal Component Analysis

```
# Check column means and standard deviations
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03

symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

The data does appear to use different units of measurement and to have slightly different variances, so scaling is necessary.

```
# Perform PCA on wisc.data by completing the following code
wisc.pr <- prcomp(wisc.data, scale. = TRUE)

summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Question 4

From your results, what proportion of the original variance is captured by the first principal components (PC1)?

```
# Using wisc.pr:  
  
#find proportion of the original variance  
prop_var <- wisc.pr$sdev^2 / sum(wisc.pr$sdev^2)  
  
#find which `prop_var` is represented by PC1 (1)  
prop_var[1]
```

```
[1] 0.4427203
```

The first principal components capture ~44% of the original variance.

Question 5

How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

```
# Using `prop_var` from above:  
  
# find cumulative variance  
cum_var <- cumsum(prop_var)  
  
# find the number of PCs required to reach at least 70% variance  
which(cum_var >= 0.70)[1]
```

```
[1] 3
```

Three principal components are required to describe at least 70% of the data's.

Question 6

How many principal components (PCs) are required to describe at least 90% of the original variance in the data?



```
# find the number of PCs required to reach at least 90% variance
which(cum_var >= 0.90)[1]
```

```
[1] 7
```

Seven principal components are required to describe at least 90% of the data's variance.

Now, I will create a biplot of the `wisc.pr` using the `biplot()` function.

```
biplot(wisc.pr)
```




Question 7

What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot seems unorganized to me. The data is hard to interpret because you cannot read the labels for each plot, and all the labels crowd the data so you can't see the datapoints- making it difficult to understand. The plot also has no units.

To start fixing this:

```
# Scatter plot observations by components 1 and 2
plot(wisc.pr$x , col = diagnosis ,
     xlab = "PC1", ylab = "PC2")
```

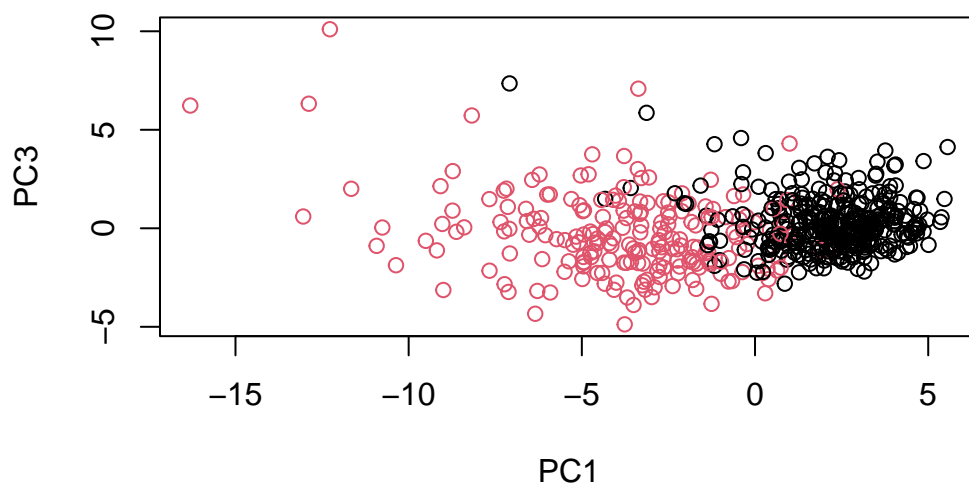


Question 8

Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

For the plots below: I notice that the new plots have slightly less clear separation when comparing them to those created by comparing PC1 vs PC2. The original plots have better separation between the data because they cover more of the variance.

```
# Repeat for components 1 and 3
plot(wisc.pr$x[, c(1, 3) ], col = diagnosis ,
     xlab = "PC1", ylab = "PC3")
```




The separation between the two groups appears to be less clear in the above graph. This is because PC2 explains more variance than PC3.

Using ggplot to make it fancier:

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col= diagnosis ) +
  labs(x="PC1 (44.27%)", y="PC2 (18.97%)") +
  geom_point()
```



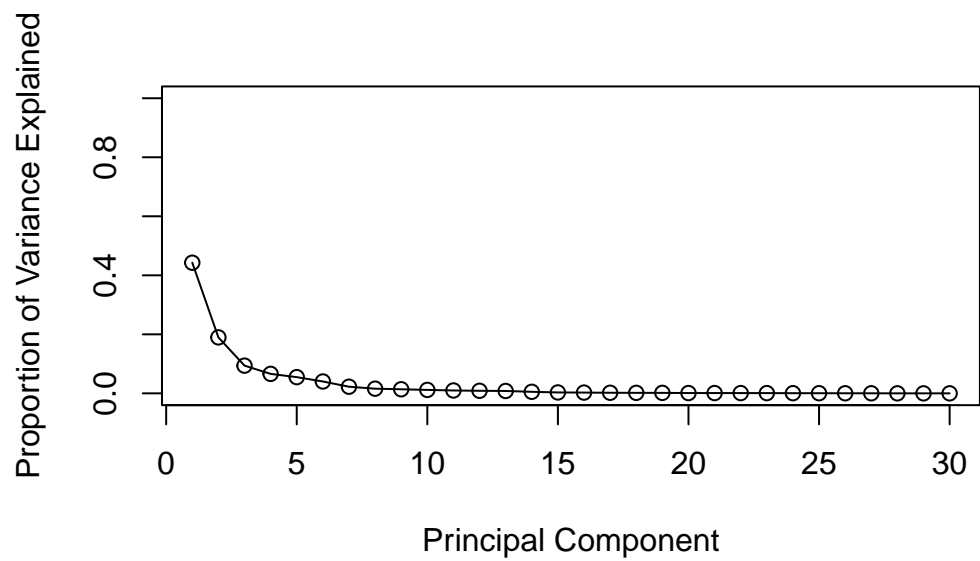
```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```


Calculate the variance explained by each principal component by dividing by the total variance explained by all principal components. Assign this to a variable called pve and create a plot of variance explained for each principal component.

```
# Variance explained by each principal component: pve
pve <- wisc.pr$sdev^2 / sum(wisc.pr$sdev^2)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



skipped the cran example

Question 9

For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
# Get the component of the loading vector for concave.points_mean (PC1)

loading_concave_points_mean <- wisc.pr$rotation["concave.points_mean", 1]

loading_concave_points_mean
```

```
[1] -0.2608538
```

The component of the loading vector for `concave.points_mean` is -0.26.

Question 10

What is the minimum number of principal components required to explain 80% of the variance of the data?

```
which(cum_var >= 0.80)[1]
```

```
[1] 5
```

There are 5 principal components required to explain 80% of the data's variance.

3. Hierarchical Clustering

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)

# Calculate distances between all pairs in `data.scaled`
data.dist <- dist(data.scaled)

# Use complete linkage to create a hierarchical clustering model. Specify the method argument
wisc.hclust <- hclust(data.dist, method = "complete")
```


Question 11

Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
# plot the dendrogram of the clustering model
plot(wisc.hclust)

# find where h = x has 4 clusters, done by trial and error
abline(h = 19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

The clustering model has 4 clusters when $h = 19$.

```
# use cutree to cut the tree down to 4 clusters  
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)  
  
# make a table to compare cluster membership to diagnoses  
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

When four clusters were picked, cluster 1 appears to correspond with mostly malignant cells ($M = 165$, $B = 12$) and cluster 3 largely corresponds to benign cells ($M = 40$, $B = 343$), but not perfectly.

Question 12

Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?


```
# use cutree to cut the tree down to 2 - 10 clusters, done by trial and error
wisc.hclust.clusters <- cutree(wisc.hclust, k = 5)

# make a table to compare cluster membership to diagnoses
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	0	5
3	343	40
4	2	0
5	0	2

Based on the options ($k = 2-10$) I tried, I ended up liking the arrangement when cutting the tree down to 5 clusters.

In this result, we have a relatively clear separation between benign and malignant cell populations, like when $k = 4$. The numbers of benign and malignant cells in clusters 1 and 3 are consistent.

However, we have the additional benefit of a more clear separation in the other three clusters (2, 4, and 5) with all three containing either only malignant or only benign cells.

So, I believe $k = 5$ presents a better diagnoses versus clusters match.

There are a number of different methods that can be used to combine points during the hierarchical clustering procedure. These include “single”, “complete”, “average” and (Prof. Grant’s favorite) “ward.D2”.

Question 13

Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
# Trying hierarchical clustering with different methods

# "single"
wisc.hclust.s <- hclust(data.dist, method = "single")

wisc.hclust.s.clusters <- cutree(wisc.hclust.s, k = 5)

table(wisc.hclust.s.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.s.clusters	B	M
1	356	209
2	1	0
3	0	1
4	0	1
5	0	1

```
# "complete"
wisc.hclust.comp <- hclust(data.dist, method = "complete")

wisc.hclust.comp.clusters <- cutree(wisc.hclust.comp, k = 5)

table(wisc.hclust.comp.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.comp.clusters	B	M
1	12	165
2	0	5
3	343	40
4	2	0
5	0	2

```
# "average"
wisc.hclust.avg <- hclust(data.dist, method = "average")

wisc.hclust.avg.clusters <- cutree(wisc.hclust.avg, k = 5)

table(wisc.hclust.avg.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.avg.clusters	B	M
1	355	208
2	2	0
3	0	1
4	0	2
5	0	1

```
# "ward.D2"
wisc.hclust.ward <- hclust(data.dist, method = "ward.D2")
```

```
wisc.hclust.ward.clusters <- cutree(wisc.hclust.ward, k = 5)

table(wisc.hclust.ward.clusters, diagnosis)
```

		diagnosis	
wisc.hclust.ward.clusters		B	M
1	0	59	
2	0	56	
3	6	48	
4	337	48	
5	14	1	

I decided to use $k = 5$ for all my comparisons, because it was what I liked most in the original method. I ended up liking “ward.D2” the most, because I believe it displayed a more thorough separation, with potentially another variable clustering them. There are several clusters that have only benign or malignant cells, with larger populations in each of the clusters than the other methods, which I just like more, visually.

4. K-means clustering (skipped)

Question 14 (skipped)

5. Combining Methods

Using the minimum number of principal components required to describe at least 90% of the variability in the data, create a hierarchical clustering model with the linkage method = “ward.D2”. (Ward’s criterion is used here because it is based on multidimensional variance like principal components analysis)

Assign the results to wisc.pr.hclust.

```
# Obtain cumulative proportion
cum_var <- cumsum(pve)

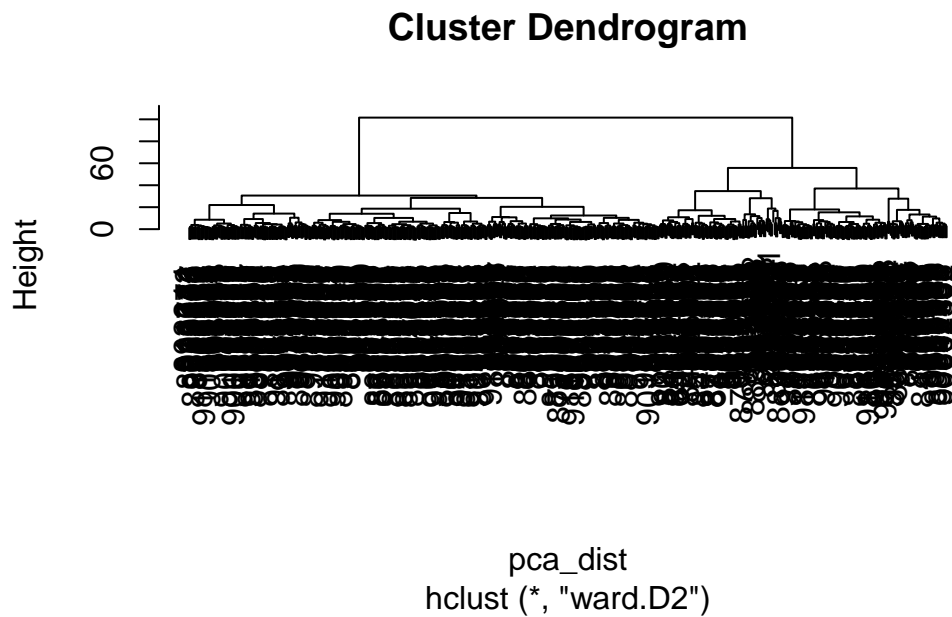
# Find number of PCs needed to reach 90%
num_pcs_90 <- which(cum_var >= 0.90)[1]

# Subset PCA-transformed data to those components
pca_subset <- wisc.pr$x[, 1:num_pcs_90]
```

```
# Compute distances using Euclidean metric on PCA scores
pca_dist <- dist(pca_subset)

# Perform hierarchical clustering with Ward's method
wisc.pr.hclust <- hclust(pca_dist, method = "ward.D2")

#Plot the new subset
plot(wisc.pr.hclust)
```



What do the two main branches of dendrogram (indicating two main clusters) show? Are the malignant and benign?


```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
  1  2
216 353
```

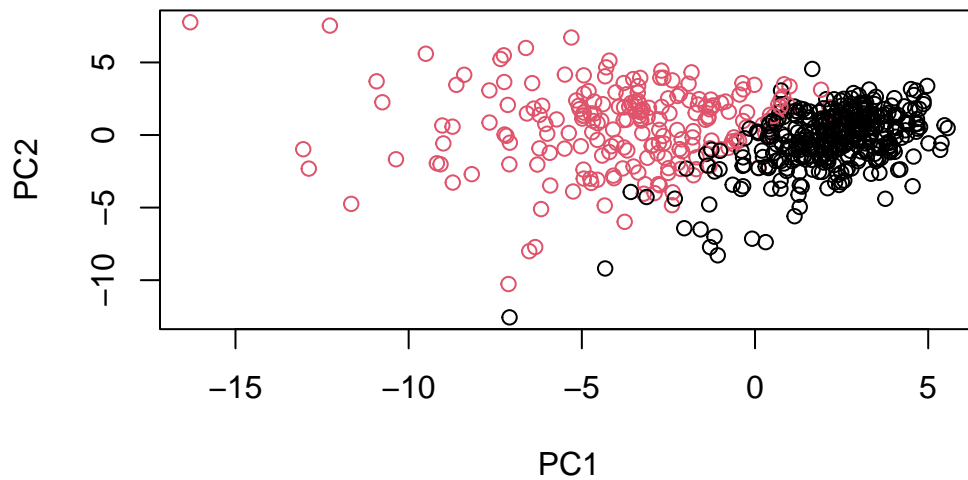
```
table(grps , diagnosis)
```

```
diagnosis
grps  B  M
1    28 188
2   329  24
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



There is a color swap here indicating the hclust cluster 1 is mostly “M” and cluster 2 is mostly “B” as we saw from the results of calling `table(grps, diagnosis)`.

To match things up we can turn our groups into a factor and reorder the levels so cluster 2 comes first and thus gets the first color (black) and cluster 1 gets the second color (red).


```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
# Plot using our re-ordered factor
plot(wisc.pr$x[,1:2], col=g)
```



I skipped the rest of the fancy stuff...

Question 15

How well does the newly created model with four clusters separate out the two diagnoses?

```
# Compare to actual diagnoses
grps <- cutree(wisc.pr.hclust, k=2)

table(grps , diagnosis)
```

```
      diagnosis
grps   B    M
1    28 188
2   329  24
```

It separates them relatively well, but not perfectly as it has grouped a small population of benign cells with the larger population of malignant cells, and vice versa.

Question 16

How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
# No table created for wisc.km$cluster because that section was labeled "Optional" and I did  
  
# Create table for `wisc.hclust.clusters`  
table(wisc.hclust.clusters , diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	0	5
3	343	40
4	2	0
5	0	2

According to my classmates (source: Grace), the k-means model did not do well. I think the hierarchical clustering model does a pretty decent job.

6. Sensitivity / Specificity

Sensitivity: refers to a test's ability to correctly detect ill patients who do have the condition. In our example, the sensitivity = the total number of samples in the cluster identified as predominantly malignant (TP) (cancerous) divided by the total number of known malignant samples (TP+FN). In other words: $TP/(TP+FN)$.

Specificity: relates to a test's ability to correctly reject healthy patients without a condition. In our example, specificity is the proportion of benign (not cancerous) samples in the cluster identified as predominantly benign that are known to be benign. In other words: $TN/(TN+FN)$.


Question 17 (skipped, running out of time)

7. Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Question 18

Which of these new patients should we prioritize for follow up based on your results?

According to the data above, the largely benign cells should be black and the largely malignant cells should be red. Therefore, because we aim to prioritize patients that are more at risk for malignant cells (red), we should prioritize patient 2 because they are located in a cluster of largely malignant cells.